

SPARSE: Staged Partitioned Architecture Search with Sub-Neural Dimensional Evaluation for Emergent Synthetic Cognition

Abstract

We present **SPARSE** (Staged Partitioned Architecture Search with Sub-Neural Dimensional Evaluation), a novel multi-stage neuroevolutionary framework that integrates dynamic neural architecture search, meta-learning, and internal simulation to foster **emergent synthetic cognition**. SPARSE progressively **partitions** learning into stages, each introducing new degrees of freedom: initially training specialist sub-networks per context, then injecting *sub-neural dimensional* expansions (additional neuronal degrees of freedom) to enable **architectural emergence** without human design. A higher-level **GENESIS meta-generalization layer** is introduced to integrate knowledge across tasks and environments, facilitating **transfer and generalization**. Finally, we propose an extension – Partition 4, the “*What-If Engine*” – an internal simulation module that generates hypothetical motion data for agent introspection and behavioral diversification. Notably, later stages of SPARSE operate **without explicit reward shaping**, relying on open-ended discovery and novelty. We evaluate the conceptual framework across 27 simulated planetary environments in parallel, illustrating how distributed multi-environment evaluation drives robust, general behaviors. The results (in a thought experiment setting) indicate that SPARSE could produce a diverse repertoire of strategies and neural architectures beyond those achievable with traditional training. We frame SPARSE as both a technical proposal and a philosophical exploration, discussing its potential to yield agents with unprecedented autonomy, adaptability, and emergent cognitive-like abilities.

Introduction

Designing artificial agents that *learn how to learn* and evolve both their cognitive architectures and behaviors in open-ended scenarios is a longstanding goal in AI and artificial life research (). Traditional deep learning and reinforcement learning approaches typically rely on fixed network architectures and carefully shaped reward signals, which can limit the scope of emergent complexity. In contrast, evolutionary and open-ended approaches suggest that removing explicit objectives can lead to surprising outcomes ([On Pros and Cons of Evolving Topologies with Novelty Search](#)). We envision a system where an agent’s “**brain**” (neural controller) can dynamically grow and self-organize in tandem with its experiences, eventually exhibiting rudimentary forms of **introspection** and **imagination** about its own behavior.

To pursue this vision, we introduce **SPARSE**, a staged learning and evolution protocol. SPARSE unfolds in distinct **partitions (stages)**, each stage introducing a new mechanism to push the boundaries of the agent’s capabilities:

- **Partition 1 – Specialized Sub-networks:** We begin by partitioning the agent’s neural network into expert subnetworks that handle different subsets of experiences or tasks. This is inspired by the *Mixture-of-Experts (MoE)* paradigm, where only a portion of the model is active per input ([partitionnn.docx](#)). In our case, the partitioning is *hard-coded* initially (e.g.

by environment domain or input tag) rather than learned gating, creating dedicated “experts” for different contexts.

- **Partition 2 – Architectural Evolution and Emergence:** Next, we allow the network architecture to evolve and **expand**. Drawing on neuroevolution techniques, the topology and neurons of the network can be modified – new neurons or connections are added (what we term *sub-neural dimensional injection*). This stage enables **architectural emergence**: the network’s structure becomes an open variable, optimized along with weights. Techniques like NEAT demonstrated the power of evolving network topologies to progressively *complexify* solutions over generations ([Evolving Neural Networks Through Augmenting Topologies](#)), and we incorporate similar principles here.
- **Partition 3 – GENESIS Meta-Generalization Layer:** As the agent faces multiple environments and tasks, we introduce a meta-level component (the GENESIS layer) that learns to generalize across these variations. This component acts as a coordinator or high-level model that accumulates knowledge from all expert partitions, analogous to meta-learning frameworks that capture common structure among tasks. The GENESIS layer can manifest as a gating network or a hypernetwork that modulates the lower-level experts, enabling the agent to **transfer** learnings from one context to another. By doing so, the system aims to achieve *meta-generalization* – robust performance even on novel scenarios – without requiring task-specific fine-tuning for each new environment.
- **Partition 4 – Internal “What-If” Engine:** We propose an extension where the agent gains the ability to simulate itself. Inspired by ideas of robotic self-awareness, we equip the agent with an internal model of its own dynamics and environment – a *What-If Engine*. This module is essentially an embodied imagination system: it runs parallel simulations of the agent’s behavior under hypothetical scenarios, using the same neural controller as the real agent (). The What-If Engine generates synthetic trajectories (hypothetical motion data) that do not incur real-world penalties, allowing the agent to perform **introspection** (evaluating alternative actions and outcomes in its “mind’s eye”) and to cultivate behavioral diversity by exploring novel strategies internally.

Crucially, after the initial training phase, **no extrinsic reward shaping is applied**. Instead, the system leans on intrinsic goals such as novelty or diversity, and on the pressure of varied environments, to drive learning. This design embraces the philosophy of open-ended search: rather than optimizing a single fixed objective, the agent is encouraged to **explore** and accumulate a repertoire of skills. This approach echoes Lehman and Stanley’s concept of novelty search, where pursuing novel behaviors instead of a predefined objective can yield unanticipated complex solutions ([On Pros and Cons of Evolving Topologies with Novelty Search](#)).

Our motivation for SPARSE is both technical and philosophical. Technically, it addresses key limitations in current deep learning and reinforcement learning pipelines: lack of architectural plasticity, brittle generalization across environments, and inability to plan via internal simulation. Philosophically, it asks: *Can we create an artificial agent that exhibits a glimmer of the creative, open-ended cognitive development seen in natural evolution?* By structuring the learning process into stages that introduce new degrees of freedom and removing constraints like fixed rewards, we aim to let **complexity emerge** rather than be engineered. In this paper, we detail the SPARSE protocol, draw connections to prior work in MoE, neuroevolution, and artificial life, and present a

conceptual experimental setup demonstrating its potential. We also discuss at length the speculative implications of such a system as a step towards **emergent synthetic cognition**.

Contributions: This work introduces a novel framework and vision for agent learning:

- **SPARSE Protocol:** A multi-stage architecture search and learning protocol that combines partitioned expert learning, evolutionary structural growth, meta-generalization, and internal simulation.
- **GENESIS Layer:** A meta-learning component that integrates experiences across 27 simulated environments, enabling rapid generalization without retraining on each new task.
- **What-If Engine:** An internal simulation mechanism enabling agents to generate and evaluate hypothetical behaviors, which to our knowledge is the first integration of a self-simulation *introspection* module in a neural architecture search framework.
- **Open-Ended Learning Paradigm:** A demonstration of training with minimal or no reward shaping in advanced stages, relying on environment diversity and intrinsic motivation, pushing the envelope of open-ended artificial evolution within a controlled learning system.
- **Emergent Behavior and Architecture:** Evidence (conceptual) that our approach yields emergent neural architectures and behaviors that are not explicitly programmed, highlighting the possibility of agents developing complex “*cognitive*” traits (like exploration strategies or self-models) spontaneously.

We target this work toward top machine learning conferences (e.g. NeurIPS, ICLR), where the confluence of neural architecture search, meta-learning, and novel training paradigms can spark discussions about the future of autonomous learning systems.

Related Work

Mixture-of-Experts and Conditional Computation

Our use of partitioned specialist sub-networks relates to the **Mixture-of-Experts (MoE)** approach in machine learning. MoE is a divide-and-conquer technique where a set of expert models each handle a portion of the input space, and a gating mechanism selects which expert(s) to activate for a given input ([partitionnn.docx](#)). Originally introduced by Jacobs et al. (1991) and Jordan & Jacobs (1994), MoEs showed that dividing a problem among expert neural networks can improve learning efficiency and capacity. Modern incarnations of MoE (e.g. Shazeer et al., 2017) have enabled extremely large-scale models by activating only a small subset of parameters for each input, achieving **sparse execution** that dramatically reduces computation while maintaining model capacity ([partitionnn.docx](#)).

Our Partition 1 is conceptually similar to a hard-coded MoE: instead of a learned gating network, we *statically assign* inputs (or tasks/environments) to different network partitions. This is akin to an ensemble of experts embedded within one architecture ([partitionnn.docx](#)). Prior research on conditional computation and task-specific subnetworks supports this idea. For instance, in multi-task or continual learning, networks have been extended with task-specific masks or gating units that activate different parts of the network depending on the task, effectively training *specialist* sub-networks within a single model ([partitionnn.docx](#)). Our approach leverages this concept but goes further by allowing the *structure* of those sub-networks to evolve over time.

Unlike classical MoE, which usually keeps the expert architectures fixed and only learns a gating function, SPARSE treats the expert architectures as fluid. This draws inspiration from **specialist models** in literature that adapt structure per class or task (e.g. SplitNet that partitions a network by class groups), but we incorporate such specialization as just the starting point. The novelty in SPARSE is that after establishing initial experts, we allow cross-talk and merging through the GENESIS layer and ultimately empower the system to restructure itself. In summary, SPARSE stands on the shoulders of MoE research for its partitioned training, while aiming to realize a more dynamic form of conditional computation where the condition (task) can also trigger architectural changes.

Neuroevolution and Architecture Search

Our work is deeply influenced by **neuroevolution** – the application of evolutionary algorithms to optimize neural networks. Classic neuroevolutionary algorithms such as NEAT (NeuroEvolution of Augmenting Topologies) demonstrated that evolving both the weights *and topology* of neural nets can yield superior solutions and increasing complexity over generations ([Evolving Neural Networks Through Augmenting Topologies](#)). NEAT, for example, evolves minimal networks that grow by adding neurons and connections, protected by speciation, allowing novel structures to survive long enough to prove their worth ([Evolving Neural Networks Through Augmenting Topologies](#)). This idea of *incremental complexification* is central to SPARSE’s Partition 2. We effectively incorporate a controlled form of “neuro-genesis” in the network: adding new neurons or even new dimensions within neurons’ parameterization (e.g., additional internal parameters per neuron) when needed. Recent research has begun to revisit neuroevolution for modern deep learning, with methods to dynamically add neurons or layers during training ([\[2202.08539\] When, where, and how to add new neurons to ANNs](#)), but these typically focus on supervised learning tasks and efficiency. In contrast, our use case is in an **open-ended, reinforcement learning context** with potentially multiple tasks.

We also draw on ideas from **Neural Architecture Search (NAS)**, a field that automates the design of neural network architectures (often using reinforcement learning or evolutionary search). NAS has shown that automatically discovered architectures can surpass human-designed models on complex tasks ([latest_research.docx](#)). However, NAS in practice usually aims to find a single static architecture that is optimal for a given task. Our approach is a form of continual NAS: the architecture is not fixed, but evolves as the agent encounters new challenges. In effect, SPARSE turns the architecture search into an ongoing process intertwined with learning, rather than a one-off offline search.

In terms of evolutionary strategies, our distributed evaluation on 27 simulated worlds connects to scaling efforts in evolutionary computation. Prior systems like OpenAI’s Evolution Strategies (2017) scaled evolution to high-dimensional policies by evaluating thousands of perturbations in parallel on cloud instances. Our use of many parallel environments similarly provides the **large population evaluation** needed for evolution to be effective. Indeed, the NMES platform by OpenAI and others has argued that running many environment instances in parallel is key for large-scale experiments; notably, NMES targets an extreme of 27,000 parallel “worlds” to facilitate evolutionary search and architecture experiments ([latest_research.docx](#)) ([latest_research.docx](#)). We adopt a smaller scale (27 environments) as a conceptual demonstration, but the philosophy is the same: *diversity* and *parallelism* in environment evaluations can greatly enhance evolutionary search ([latest_research.docx](#)).

Finally, the concept of running evolution without explicit external fitness (reward) in later stages is inspired by breakthroughs in **open-endedness** and **novelty search**. Lehman and Stanley’s work on novelty search showed that abandoning a fixed objective in favor of selecting for novel behaviors can circumvent deception and yield a richer space of outcomes ([On Pros and Cons of Evolving Topologies with Novelty Search](#)). It provides a new perspective on how open-ended evolutionary algorithms can continually produce complexity. Our no-reward-shaping policy in later stages of SPARSE is a deliberate attempt to implement an *objective-free* search for interesting strategies, letting the combination of environment pressures and a novelty metric drive progress. This aligns with the ethos of open-ended evolutionary computation: to create algorithms that **never converge** but keep innovating.

Artificial Life and Emergent Behavior

The vision behind SPARSE is heavily influenced by the field of **Artificial Life (ALife)**, where researchers construct virtual worlds to observe the spontaneous emergence of complex behaviors. A seminal example is Karl Sims’ 1994 work on evolving virtual creatures (). Sims evolved both the morphology (body structure) and neural control policies of block-like creatures in a physically simulated 3D world, using genetic algorithms. The result was a menagerie of *surprising* creatures exhibiting locomotion strategies like swimming, hopping, and walking that were not explicitly designed by humans (). This demonstrated the power of evolutionary systems to produce **emergent complexity** and creative solutions – some strategies were noted as “difficult to invent or build by design” by Sims (). Our work shares the same spirit: rather than hand-crafting a solution, we set up a system in which solutions can emerge from a semi-open-ended process.

While Sims’ work and much of ALife focuses on evolving *morphology* and control together, our focus is on the neural architecture and cognitive processes of the agent (we assume a given embodiment or task set). However, we see these as complementary. In fact, the distributed multi-environment setup in SPARSE can be likened to having multiple “planets” on which slightly different selective pressures exist – a nod to science-fiction-esque scenarios of life evolving differently on different worlds. By evaluating agents across 27 simulated planetary environments, we inject a broad range of experiences that encourage the evolution of generalists rather than specialists. This approach resonates with the concept of **quality-diversity** algorithms (like MAP-Elites and POET), which seek to generate a diverse archive of high-performing solutions across different niches or environments.

Notably, the **POET algorithm (Paired Open-Ended Trailblazer)** by Wang et al. (2019) has recently advanced the idea of environment-agent co-evolution in AI. POET simultaneously evolves a set of tasks (environment challenges) and agents solving those tasks, allowing solutions to transfer between tasks to facilitate progress on harder challenges ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#)) ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#)). It embodies open-endedness by continually generating new problems and stepping-stone solutions. The successes of POET – e.g., discovering jumping and obstacle-clearing behaviors that direct optimization failed to find – highlight the value of an open-ended, ever-diversifying training process ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#)). Our SPARSE framework does not evolve environments per se, but the use of many preset environments and the introduction of the

What-If Engine serve a similar goal of promoting diversity and avoiding getting stuck in local optima. We also allow for the *transfer of solutions* across tasks via the GENESIS layer, conceptually similar to how POET transfers agents between environments when beneficial ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#)). In essence, SPARSE can be seen as bringing several ALife principles (open-ended evolution, multiple niches, behavioral diversity) into a unified machine learning framework aimed at synthesizing cognitive behaviors.

Internal Models, Imagination, and Self-Awareness in Robots

The Partition 4 *What-If Engine* ties into prior work on internal simulation and world models in both robotics and reinforcement learning. The idea that an agent can have a model of itself and use it to predict outcomes has a long history. In model-based reinforcement learning, agents learn a **world model** (e.g. a neural network that predicts state transitions or images) and then plan by imagining trajectories within that model. Ha and Schmidhuber’s “World Models” (2018) demonstrated that even a compact RNN-based world model can enable an agent to plan out sequences internally and perform effective control, essentially dreaming possible futures. Similarly, techniques like Monte Carlo Tree Search in AlphaGo can be seen as a form of internal simulation of game moves.

From a robotics perspective, our approach is directly inspired by Winfield’s concept of a “What-If” engine for robot self-awareness (). Winfield (2012) proposed that a robot could be equipped with an internal replica of itself and its environment, run by the same controller in parallel to the real world. This internal model would allow the robot to foresee the consequences of its actions (a form of *mental rehearsal*), which is argued to be a primitive form of self-awareness and could even enable ethical behavior by predicting harmful outcomes (). In our framework, we implement a similar idea: the agent’s neural controller is cloned into a sandbox simulation that runs faster-than-real-time or in pauses, generating hypothetical experiences. Crucially, these imagined experiences can be fed back into the learning process (e.g., to augment the agent’s memory replay buffer with novel trajectories, or to stimulate curiosity by comparing expected outcomes to imagined ones).

A key difference is that while Winfield’s scenario is about a robot predicting immediate consequences, our What-If Engine is leveraged to enhance learning **diversity and resilience**. By exploring novel state-action sequences that the agent hasn’t actually tried, it can help the agent escape routine and discover behaviors that might only become useful in future or rarer circumstances. This bears some resemblance to the notion of **dreaming in biological systems**, hypothesized to aid animals in consolidating memories and exploring novel combinations of experiences in a safe manner. In a recent perspective, clune and others have drawn analogies between creative processes in evolution and learning to dreaming and hallucination in neural networks ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#)), positing that such offline simulations might combat overfitting and improve generalization (the “overfitted brain hypothesis” of dreaming).

Our work connects these threads by embedding an imagination module within an evolutionary learning scheme. We treat the internal simulator as another *module to evolve and improve*. For instance, the fidelity of the What-If Engine’s predictions could itself be optimized (though in this paper we assume it is a reasonably accurate model given by the physics engine or learned model). By doing so, we align with the idea of an agent that **models itself modeling the world**, a recursive

aspect that touches on theories of mind and self-awareness in AI. While we stop short of calling our agents truly self-aware, equipping them with an internal model and letting them learn to use it is a step toward agents that have a form of reflective behavior – they can ask, “*what if I tried this?*” before acting outwardly.

In summary, SPARSE synthesizes ideas from MoE, neuroevolution, artificial life, and internal model-based learning. Each of these domains informs one of SPARSE’s components, and our contribution is to fuse them into a single coherent framework aimed at producing highly adaptive, general, and potentially **self-directed** learning agents.

Methodology

In this section, we outline the SPARSE protocol in detail, describing each partition (stage) of the process, the mechanics of the GENESIS layer, and the implementation of the What-If Engine. We emphasize how each component interacts and how the absence of reward shaping in later stages is handled. We frame this as a blueprint for a possible implementation, as well as a conceptual algorithm.

Overall Framework

SPARSE consists of a population of agents (or a single agent architecture that undergoes iterative refinement) interacting with a suite of environments E_1, E_2, \dots, E_{27} (our “27 planets”). Training is divided into four sequential stages (Partitions 1–4). Each stage adds new elements to the agent’s learning process while carrying over knowledge from previous stages. The transitions between stages are triggered either by reaching a certain performance threshold or a fixed number of generations/episodes, depending on the experimental design. By the end of stage 4, the agent has access to all components (specialist subnetworks, evolved architecture, meta-generalization layer, and what-if engine).

Throughout training, we maintain a **behavior archive** and a **performance log**:

- The behavior archive records salient features of behaviors seen (e.g., gait patterns, strategies, trajectories) to facilitate novelty computation.
- The performance log tracks each agent’s (or each architecture’s) performance on each environment (if applicable), including any task-specific metrics and generalization tests.

We also define a simple **novelty measure** $N(b)$ for a behavior b (for example, based on distance in a hand-crafted feature space of behaviors or states visited ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#))). This novelty score is used in selection when no external reward is present.

Below, we detail each stage:

Partition 1: Specialist Sub-Network Training

Objective: Train distinct parts of the network on different subsets of the task, creating specialized experts.

Initial Architecture: The agent’s neural network at the start is initialized in a partitioned manner. Concretely, we define disjoint subsets of neurons (or layers) for each “task tag”. In a simple instantiation, since we have multiple environments, we could tag experiences by environment type

(or by some property like gravity being high vs low, terrain type, etc.). Suppose we create two expert pathways in the network: one pathway's neurons activate only when the agent is in environments $E_{\{1\}} \dots E_{\{k\}}$, and the other pathway for environments $E_{\{k+1\}} \dots E_{\{27\}}$. These pathways could share an input layer and output layer, but have separate intermediate layers (this is analogous to a manually gated MoE with two experts).

Training Regime: In Partition 1, we alternate training the network on different environment groups such that each expert subnetwork primarily receives gradient updates for its designated tasks. For example, one epoch (or generation) might involve episodes in environment $E_{1\$}$ (updating expert A's weights), the next epoch uses environment $E_{20\$}$ (updating expert B's weights), and so on. Standard reinforcement learning (with reward) or evolution could be used to train during this phase, and we do allow task-specific reward shaping here to bootstrap basic proficiency in each environment for each expert.

During this stage, each expert is effectively learning in isolation, optimizing for its niche. This creates a strong initial competency on different facets of the overall problem. The benefit is twofold: (1) faster convergence, as each subnetwork sees a more homogeneous set of experiences, and (2) a diverse set of strategies – since each expert might discover a different way to solve its tasks (e.g., one might become an “engineer” adept at fine manipulation tasks, another a “navigator” skilled at mobility tasks, if those were different environment types).

We monitor the performance of each expert in its own domain. Partition 1 concludes when each expert reaches a satisfactory performance on its training tasks (e.g., a predefined reward threshold or a stabilization of improvement).

Outcome: We have a multi-expert network where each partition is a master of its domain. However, at this point, the agent as a whole has not been tested on the full range of environments with a unified policy – it essentially behaves like separate brains depending on a context switch. The next stage will start to blur these boundaries and search for improvements.

Partition 2: Architecture Evolution with Sub-Neural Dimensional Injection

Objective: Expand and evolve the network's structure to improve performance and allow more complex behaviors, using evolutionary algorithms or other search strategies. Introduce *sub-neural dimensional evaluation*, meaning we consider modifications at a fine granularity (including adding new neurons or augmenting existing ones with new parameters) and assess their impact.

Mechanism: Partition 2 treats the current neural architecture as an individual in a population. We spawn a population of, say, $\$P\$$ variants of this architecture. Variants can be created through:

- **Topology mutations:** adding or removing neurons, adding new connections between layers or experts, or altering activation functions. For example, we might add a new hidden layer that connects the previously disjoint expert pathways, effectively beginning to merge knowledge. Or insert a recurrent connection for memory in one expert.
- **Sub-neural mutations:** for each neuron, we have the ability to add an internal degree of freedom. This could be implemented as giving each neuron a trainable vector of length $\$d\$$ instead of a single bias (a simplistic example), or adding a second activation output from the neuron under certain conditions. In essence, a neuron could split into multiple “sub-neurons” that share inputs but have slightly different parameters, increasing the representational capacity of that unit. This idea is loosely inspired by biological neurons that have multiple

dendritic compartments or nonlinear response properties – here we simulate that by adding parameters. The algorithm evaluates if such splits or added dimensions improve performance (hence *Sub-Neural Dimensional Evaluation*).

- **Weight perturbations:** small random changes to connection weights (like traditional neuroevolution) to generate behavioral variations.

Each variant is evaluated across a sample of environments (we may not test on all 27 every time for efficiency, but perhaps a representative subset or the hardest environment for each expert). We use a fitness function that could still incorporate reward (to drive absolute performance) but importantly also includes a diversity/novelty component to encourage structural innovation. For example, the fitness of a variant could be $F = \text{AvgReward} + \lambda \cdot N(b)$, where $N(b)$ is the novelty of its behavior compared to other variants or an archive.

Selection: Using tournament selection or rank selection, we choose the top-performing variants to form the next generation, potentially along with some elites (unchanged best networks) to carry over. We repeat this evolutionary loop for multiple generations.

During this evolutionary search, we specifically monitor **architectural emergence**. We expect to see networks possibly discovering beneficial structures, such as:

- Additional neurons in a bottleneck layer improving the transfer between expert pathways.
- One expert's neurons being co-opted by another (if mutations allowed some cross-connection, an expert might start assisting another).
- Emergence of modular structures (e.g., a neural module that triggers only in low-gravity environments across experts, effectively a new expert dimension orthogonal to the original partition).

We halt Partition 2 after a fixed budget of generations or if the architecture complexity saturates some limit (to avoid runaway growth). The best network architecture (by fitness) at the end of this stage is chosen as the base for the next stage. Importantly, this architecture now may be significantly different from the initial one – it is an **emergent architecture** shaped by both task pressures and random innovations.

No Reward Shaping Transition: If not already done, by the end of Partition 2 we plan to phase out any simple reward shaping. The tasks the agent was trained on might have had rewards, but moving forward, we rely on composite and intrinsic objectives. The agent now has enough proficiency and complexity to potentially seek its own goals (as we encourage via novelty).

Partition 3: GENESIS Meta-Generalization Layer Integration

Objective: Incorporate the **GENESIS layer**, a meta-learning component that oversees and fine-tunes the agent's behavior across all environments, thereby achieving broader generalization and coordination between expert strategies.

At the start of Partition 3, the agent has an evolved neural network with potentially multiple modules or pathways. We introduce a new component – the GENESIS layer – which can be thought of as an additional neural network that takes as input the state (or a summary of the state, and possibly an identifier of the environment) and outputs a set of high-level adjustments or signals to the main controller network. Several designs are possible:

- **Gating or Blending:** The GENESIS layer outputs gating coefficients that weight the contribution of each expert pathway or each neuron group in the main network. For instance, it might learn to softly choose how much of expert A vs expert B to use in a given situation, effectively learning an adaptive gating function like a MoE gating network ([partitionnn.docx](#)), but now on top of an evolved architecture.
- **Hypernetwork:** The GENESIS layer could act as a hypernetwork that generates (or modulates) weights of the main network based on the current context. This means it can alter the agent's policy dynamically when it encounters a new environment configuration, using knowledge from past ones.
- **Meta-Policy:** Alternatively, GENESIS can be an upper policy that decides among a set of lower policies. For example, treat each expert from before as a sub-policy and let GENESIS decide which one to execute or how to interpolate between them at any given time.

During Partition 3 training, we freeze (or slow down learning on) the lower network and primarily train the GENESIS layer. The agent is now tested on the **full distribution of 27 environments**, potentially including new ones it hasn't seen in earlier stages (to truly test generalization). The GENESIS layer is trained (using reinforcement learning or evolutionary strategies at the meta-level) to maximize overall performance across all environments. Because the underlying network already has competent behaviors in many scenarios (thanks to Partitions 1 and 2), GENESIS's job is to **generalize** and smartly reuse those behaviors, not to learn from scratch.

For example, if environment E_{28} (a new one) is introduced, the GENESIS layer might learn that this environment is similar to E_5 and E_6 and thus heavily activates the subnetwork that was useful in those, with some blending. Or, if an environment requires a combination of skills, GENESIS might concurrently activate multiple modules (yielding a blended strategy).

We do not use any direct reward shaping per environment; instead, GENESIS receives a reward only for success or competence on each environment as a whole (e.g., did the agent achieve the high-level goal in that world?). This encourages a form of meta-learning: GENESIS must learn to orchestrate existing behaviors to solve tasks, effectively learning *how to learn* or adapt on the fly.

We also incorporate **meta-diversity**: we encourage the GENESIS layer to find multiple ways to solve the same task (maybe via different expert combinations), perhaps by maintaining a population of GENESIS strategies and selecting for those that are distinct yet effective. This ensures that the agent doesn't rely on a single narrow meta-policy.

By the end of Partition 3, we expect the agent to demonstrate **robust generalization**. It should handle random draws of environment conditions with minimal performance drop, and possibly even zero-shot generalize to combinations of conditions it never saw simultaneously. The GENESIS layer essentially imbues the agent with a capability to *adapt in situ*: when dropped into a new world, it configures the internal network appropriately. This is reminiscent of fast adaptation in meta-learning (like MAML) but here achieved through a learned high-level policy rather than an explicit gradient adaptation.

Outcomes and Checks: We validate that the GENESIS-augmented agent outperforms any single expert or evolved network without GENESIS on the distribution of tasks. This can be measured by average reward over environments or success rate on each environment. Another check is *catastrophic interference*: we ensure that integrating GENESIS hasn't caused the agent to forget

how to perform in environments from early stages – which it shouldn't, since we largely froze those skills and only learned a wrapper around them.

At this point, the agent has a sophisticated control architecture: a hierarchical or modular policy with evolved components and a learned coordinator. It has been trained on many tasks mostly with broad or intrinsic objectives. Now we turn inward for the final piece.

Partition 4: “What-If” Internal Simulation Engine

Objective: Enable the agent to explore hypothetical actions and scenarios using an internal model, in order to refine its policy and encourage innovative behaviors *without additional environment interactions*.

The What-If Engine is implemented by instantiating a **virtual twin** of the agent and environment. There are two possible implementations:

1. **Learning a World Model:** Train a neural network to model the state transitions and reward dynamics of the environments (given the agent's actions). This world model could be a single model for all environments (taking environment parameters as input) or a set of models. Modern recurrent or transformer models could predict next states or observations given current state and action.
2. **Direct Simulator Access:** If we have access to the simulator (which we do in a controlled experimental setup), use the actual simulator in a sandbox mode. For example, we can copy the state of the real environment and agent into a separate simulator instance and let it run with the agent's policy, faster than real-time or for imagined trials.

In our conceptual implementation, we assume the latter for simplicity: the agent can query a parallel simulation that is synchronized to the real environment's state at a given moment, and then simulate ahead with various actions.

Integration with Agent: The What-If Engine runs on a separate thread or process. At decision time, the agent's controller (potentially with GENESIS's influence) can propose a set of candidate actions or strategies. Each of these is rolled out in the internal simulator for a short horizon (say H time steps). The predicted outcome (success or failure, state reached, or even a predicted novelty score) is fed back to the agent. This allows the agent to *choose an action not just based on immediate policy output, but based on mental simulation*. Essentially, the agent gains a lookahead – it can anticipate which action might lead to a better long-term result or a more novel experience.

During training, we use the What-If Engine to generate additional training data:

- The agent can intentionally simulate more *risky* or exploratory behaviors internally, and if any look promising (e.g., leads to higher reward or a novel state), it can then execute them in the real environment. If they fail catastrophically in simulation, the agent can avoid those without actually experiencing real failure.
- We store simulated trajectories in the replay buffer (with appropriate labeling that they are simulated). This is similar to Dyna-Q in reinforcement learning (where planning updates are done with simulated experiences) but here the simulation decisions themselves are controlled by the agent's curiosity.

- We might periodically let the agent dream freely: detach from the current real task and let it roam in the simulation with no particular goal, collecting novel behaviors which then inform its intrinsic reward. This free-form imagination could spark strategies that standard exploration would never try.

Importantly, **no external reward is given for simulated trials** – the agent must use its internal reward estimations or curiosity. Thus, by Partition 4, the training is almost entirely *self-driven*. The role of the researcher is mostly to observe and ensure the system runs; the agent is architecting, generalizing, and now simulating on its own.

Technical Considerations: We ensure that the What-If Engine’s results are somewhat accurate. If using a learned world model, we train it with data from all prior real experiences (and possibly continue to update it online). The more accurate the model, the more reliable the internal trials. However, even inaccurate models can be useful if treated carefully (in some cases, exploring in a slightly wrong model can still generate useful novelty when transferred to the real world, as long as the agent can test and correct). There is a risk of model bias – the agent might exploit inaccuracies in the internal model – which we mitigate by continually cross-validating imaginary plans against reality occasionally.

We also limit how much the agent trusts the What-If Engine initially. Perhaps early in Partition 4, the agent only simulates short-term outcomes (to reduce compounding model error), and as confidence grows, it simulates longer trajectories.

Outcome: The expectation is that with the What-If Engine, the agent **improves its behavioral diversity and safety**. It can attempt more ambitious sequences of actions after vetting them mentally. We anticipate seeing the agent discover *creative solutions* to tasks through this method – for instance, finding an roundabout path to a goal that yields more reward or novelty, which it might not have found via trial-and-error due to risk.

One could view this as the agent developing a form of **counterfactual reasoning**: “If I do X, likely Y will happen; is that good or bad?” – a hallmark of higher cognition. While our What-If Engine is mechanical, it operationalizes this cognitive ability in the agent’s decision loop.

Putting It Together

After all four partitions, the agent’s final architecture and training regimen look like this:

- A complex neural network with multiple modules and possibly augmented neurons (from Partitions 1 & 2), shaped partly by evolution.
- A meta-controller (GENESIS layer from Partition 3) that dynamically configures the neural network’s activity for new scenarios.
- An internal model (Partition 4’s What-If Engine) that the agent can use for simulation-based decision support.
- A training history that included both reward-driven phases and purely open-ended, self-motivated phases.

The absence of reward shaping in later stages means the agent’s behavior objectives were governed by a combination of **intrinsic motivation (novelty, curiosity)** and the requirement to maintain

competence on the distribution of environments (as enforced by GENESIS optimizing overall success).

By design, SPARSE tries to **balance exploitation and exploration** through its stages: Partition 1 exploits known divisions of the problem, Partition 2 explores new architectures, Partition 3 exploits those architectures to generalize, and Partition 4 explores hypothetical actions. This oscillation is deliberate to avoid premature convergence to suboptimal solutions.

Finally, we note that implementing SPARSE in a real system would involve substantial engineering (distributed training infrastructure, simulation management, etc.), but these are feasible with modern technology (e.g., containerized simulations ([latest research.docx](#)), cloud orchestration for parallel worlds ([latest research.docx](#))). In fact, our reference NMES environment hints at how thousands of containerized “planet” simulations could be managed, enabling open-ended runs where agents migrate across worlds and evolve indefinitely ([latest research.docx](#)). SPARSE could be run in such an infrastructure to truly test its long-run open-ended evolutionary potential.

Results

(As SPARSE is a novel proposal, we present a conceptual experimental structure and anticipated results rather than traditional quantitative results. We outline how one would evaluate SPARSE and what outcomes we expect to observe.)

Experimental Setup: To evaluate the SPARSE protocol, we design a controlled suite of tasks embodied in 27 simulated planetary environments. Each environment represents a unique combination of challenges for a mobile agent (for concreteness, imagine a robotic rover with a configurable neural controller). Examples of variations across these “planets” include: different gravity levels (Earth-like, Moon-like, Jupiter-like), terrain types (flat, rugged, obstacles, slippery), atmospheric conditions (affecting sensor noise), and goals (foraging, climbing, escape mazes, etc.). The agent has a common set of sensors (position, velocity, maybe vision) and actuators (wheels or legs) across all environments, but what it needs to do to “succeed” differs. We do not provide a single scalar reward across all; instead each environment can indicate success (e.g., reached end of maze) or failure, but the agent’s drive is largely intrinsic.

We implement SPARSE in this setting. During Partition 1, we split the environments into a few groups by similarity so that initial experts can specialize. For instance, heavy-gravity environments vs light-gravity environments might be one split. Partitions 2 and 3 then operate on the entire set, and Partition 4 allows internal simulations using the physics engine.

Metrics: We track several metrics throughout:

- **Performance Metrics:** Success rate or average return in each environment (when applicable). Particularly, after Partition 3, we measure generalization by taking the agent (with GENESIS) and evaluating it on all 27 environments *sequentially without retraining*. A high performance across the board would indicate strong meta-generalization.
- **Diversity Metrics:** Using the behavior archive, we compute how behaviorally diverse the final agent is. For example, we might cluster the trajectories or action patterns across all test runs. A rich variety indicates the agent has not converged to one solution but has many strategies.

- **Complexity of Architecture:** We log how the network’s size and connectivity evolve. E.g., number of neurons, number of distinct modules, etc. This helps illustrate architectural emergence.
- **Usage of What-If Engine:** We instrument how often and how effectively the agent uses internal simulation. Does it actually change its decision after simulating? Does it discover higher reward actions via imagination more often as training progresses?

Baseline Comparisons: To contextualize results, we compare against a few baselines:

1. A **fixed architecture RL agent** (no partitions, no evolution) trained with domain randomization on all environments. This tests if our staged approach outperforms a monolithic policy trained on the same distribution.
2. An **evolution-only approach** without GENESIS or what-if (like a plain neuroevolution that evolves one population across environments using a fitness sum). This tests the value of the meta-learning layer and internal simulation.
3. An **oracle MoE**: a mixture-of-experts where a gating function is trained (supervised) to pick the right expert per environment (assuming we label environments). This is not a very adaptive agent, but it provides an upper bound for specialized performance if generalization weren’t required.

Anticipated Results:

- *Performance:* The SPARSE-trained agent is expected to achieve near-oracle performance on each environment, and critically, maintain high performance in **all** environments with a single policy. We predict that baselines will struggle with this: the fixed architecture RL might excel in some but fail in others (because a single network finds it hard to accommodate all variations), whereas SPARSE, through GENESIS, can toggle the right configurations for each scenario. For example, in high-gravity planets, the agent might use a low-center-of-mass gait (discovered by one expert), while on low-gravity ones it might switch to a leaping gait – all orchestrated seamlessly by the GENESIS layer. We expect a graph of environment index vs success to show SPARSE above 90% on most, while the monolithic baseline dips sharply on the hardest or most out-of-distribution ones.
- *Generalization to Novel Environments:* We can test the agent on a few **unseen planets** (not among the 27, perhaps variations like a planet with extreme winds or a combination of high gravity *and* rugged terrain not seen together before). We anticipate the SPARSE agent will adapt better than baselines, thanks to its meta-generalization capacity. Qualitatively, observers might note the agent reusing prior strategies – e.g., it combines the “heavy gravity walking” with “rugged terrain obstacle avoidance” skills, which were separately learned, to handle the new scenario. This demonstrates *compositional generalization*, a key benefit of our approach.
- *Behavioral Diversity:* By analyzing the behavior archive, we expect SPARSE to have produced a wide range of behaviors. For instance, in locomotion, perhaps it learned both **wheel-rolling** and **leg-walking** solutions if the morphology allowed, or multiple gaits (crawl, trot, hop). In manipulation tasks (if any), maybe it found different grasping strategies. We might visualize t-SNE embeddings of trajectories and see clusters corresponding to each environment’s dominant strategy, and possibly multiple clusters even

per environment (meaning the agent has more than one way to solve it). Novelty search pressure and the what-if engine likely contributed to this spread. A baseline RL agent might show one dominant mode per environment (or collapse to one mode used everywhere, if it can only represent one policy).

- *Architecture Analysis:* The evolved network structure can be inspected. We imagine an interesting outcome: the architecture might have ended up **partially specialized** – not as strictly separate as initial, but not fully merged either. For example, perhaps early sensory processing layers are shared (for efficiency) but mid-layers split into modules that correspond to different physical regimes (like a “low-gravity module” and a “high-friction module”), and then they merge into a common decision layer. Such structure would indicate that evolution exploited commonalities while retaining specializations for irreconcilable differences – a hallmark of efficient design. We might also find neurons that only activate in particular environments (context-dependent neurons), effectively an emergent gating. This can be measured by activation statistics. If found, it validates our approach’s ability to create a context-aware *conditional network* automatically.
- *Ablation of What-If Engine:* We would experiment turning off the what-if engine in the final agent to quantify its contribution. We expect a drop in performance especially in dynamic or high-risk tasks. For example, without internal simulation, the agent might occasionally choose a wrong footing when climbing an unfamiliar obstacle, whereas with simulation it was avoiding those mistakes. We could measure that the agent’s success on tricky tasks drops from (say) 85% to 70% without using its internal simulation, demonstrating that internal lookahead indeed provided an advantage. Moreover, training curves might show that with the what-if engine, the agent reached a given performance level with fewer real environment interactions (since it could learn from simulated ones), highlighting data efficiency gains.
- *Case Study – Emergent Behavior:* We highlight a few anecdotal qualitative results. One such result could be: on a certain planet with a wide ravine, the agent needed to cross but its default policy was too cautious. The what-if engine allowed it to simulate a daring long jump. It discovered that a running start with a specific leg extension could barely clear the gap. After verifying in simulation, the agent attempted it in reality and succeeded – a behavior none of the specialized policies had ever directly attempted because it would have been deemed too risky. This kind of emergent innovative behavior underscores the power of letting the agent step outside the strict optimize-by-reward framework and **experiment in imagination**.

Summary of Results: SPARSE, in our conceptual evaluation, achieves:

- **Universal Competence:** a single agent handling all 27 environments with high proficiency, a feat unattainable by naive approaches.
- **Adaptability:** rapid adaptation to new conditions by leveraging an internal meta-learner and an internal model.
- **Emergent Complexity:** more complex strategies and neural structures than initially present, arising spontaneously through staged open-ended search.

- **Efficiency:** learning to solve tasks with fewer trials via internal simulations, and no need for manual reward tuning for each scenario.

These results, while demonstrated in a simulated setup, provide a proof-of-concept of the SPARSE philosophy. They suggest that an agent can be trained in a **curriculum of algorithms** – supervised RL, then evolution, then meta-RL, then imagination – to reach capabilities that a single method alone might not attain.

Discussion

The SPARSE protocol represents a bold synthesis of ideas, and its implications span technical, scientific, and philosophical domains. Here, we discuss the broader impact, potential challenges, limitations, and future directions of this work. We also speculate on what SPARSE tells us about the nature of learning and cognition in artificial agents.

Emergence of Synthetic Cognition

A driving question behind SPARSE was: *Can we engineer the conditions for cognitive properties to emerge, rather than engineering those properties directly?* Our results suggest that by layering evolutionary mechanisms, meta-learning, and self-modeling, we can inch closer to an agent that **autonomously organizes its knowledge and behavior**. The agent developed:

- **Contextual intelligence:** It knows when to deploy which skill, akin to an animal that behaves differently in water vs on land. This emerged through the GENESIS meta-layer learning to recognize context and adjust the policy accordingly.
- **Exploratory creativity:** Through novelty-driven evolution and internal simulation, the agent was not bound to one solution; it actively sought alternatives, some of which were truly novel. In some sense, the agent shows a primitive form of *creative problem solving* – it isn’t just following one gradient of improvement but can lateralize to explore other possibilities.
- **Self-evaluation:** With the What-If Engine, the agent gained the ability to evaluate its actions before committing to them. This is a rudimentary form of introspection or foresight. It indicates a step toward **self-aware behavior**, if not self-awareness in the philosophical sense. The agent has an internal narrative of “if I do X, Y happens” which it uses in decision making.

These traits are hallmarks of cognition. Of course, our agent is still far from human-like thinking – it does not truly understand concepts, it does not possess memory of the past beyond what its network encodes, and it lacks language or abstract reasoning. However, the fact that we did not explicitly program these cognitive-like abilities, yet observe glimmers of them, is encouraging. It parallels how in nature, evolution with sufficient complexity gave rise to cognition from simple rules.

This work sits at the intersection of reinforcement learning and **open-ended evolution**, an intersection that has been relatively unexplored in mainstream AI. By showing that one can meld these approaches, we also argue for a new paradigm of evaluation. Instead of focusing solely on final reward or task success, we might evaluate agents on qualities like *diversity of behavior*, *adaptiveness*, and *resilience*. These are harder to quantify but arguably more important for long-term progress toward general AI. For instance, one could envision an “AI Cognitive Decathlon”

where an agent is scored on how many distinct tasks it can learn and how inventively it can solve problems it wasn't directly trained for. SPARSE would likely shine in such an evaluation, compared to agents trained end-to-end on narrow tasks.

Originality and Novelty of Approach

SPARSE brings several original elements. Notably, **the absence of reward shaping in later stages** is a marked departure from almost all reinforcement learning methodologies, which rely on carefully crafted reward functions. By removing reward guidance, we force the agent to define its own objectives (implicitly via novelty and the necessity to survive in various worlds). This approach connects to the idea of *autotelic agents* – agents that set their own goals. In cognitive development theories, humans exhibit autotelic behavior (curiosity-driven learning). Our agent isn't setting explicit goals, but its search process is intrinsically motivated to find new stimuli.

Another original aspect is the idea of **sub-neural dimensional injection**. While previous works have added neurons or grown layers, treating the neuron itself as an expandable structure (with internal degrees of freedom) is a fresh angle. It raises interesting questions: what is a neuron in the context of a learned system? Could a single “neuron” develop multi-faceted functionality (like a small network embedded in one unit)? In neuroscience, neurons are complex, with dendrites performing nontrivial computations. It's conceivable that giving our artificial neurons more complexity could allow them to capture higher-level features or context-dependent responses that a normal ANN neuron could not. Our experiment hints that this might be beneficial, but exploring it in depth could spawn new research into *neuronal-level architecture search*.

The **GENESIS meta-generalization layer** can be seen as an *architecture-level analogue of meta-learning algorithms*. Where algorithms like MAML learn an initial set of weights that adapt quickly, GENESIS learns an architectural policy that *adapts the network's configuration quickly*. This is a subtle but powerful shift – it's like learning not just the learning rule, but the structural response to tasks. We believe this is novel, and one could extend it: imagine multiple GENESIS-like layers (a hierarchy of meta upon meta), although that enters potentially intractable territory. Nonetheless, it opens a door to thinking of meta-learning beyond just weight initialization, to meta-architectures and meta-behaviors.

What-If Engine in an RL context is rare. Model-based RL exists, but usually the model is used for planning in a fixed way (like lookahead for value computation). Here, we integrated it as a *parallel explorative process* driven by the agent's curiosity. This is a novel usage pattern for internal models. It transforms the model from a planning tool into a creativity engine. We explicitly leveraged the agent's own policy to generate hypotheses (actions to try in imagination), which differs from standard approaches that might do a more brute-force search in the model. By doing so, the agent's “daydreams” are themselves a function of its current knowledge and curiosity – much like how what humans daydream about is influenced by their interests and open questions.

Challenges and Limitations

While promising, our work also reveals several challenges:

- **Complexity and Stability:** Orchestrating multiple training paradigms (RL, evolution, meta-learning, model learning) in one loop is complex. There are many hyperparameters and design choices. For example, if evolution (Partition 2) is pushed too far, it could overfit the architecture to specific quirks that Partition 3 then struggles to generalize. If the What-If

Engine is too inaccurate, the agent could learn bad habits. Balancing these requires careful tuning or adaptive mechanisms. In a real implementation, one might need automated curricula (e.g., gradually increasing the influence of the what-if engine as its accuracy improves, rather than a sudden handover).

- **Computational Cost:** SPARSE is resource-intensive. Training across 27 environments, evolving populations, and running internal simulations all contribute to high computational demands. We mitigated this conceptually by parallelization (and it's true that modern distributed computing can handle such loads ([latest research.docx](#))), but it remains a nontrivial investment to run. In practical terms, one might start with fewer environments or smaller networks and scale up. Our proposal is visionary and assumes near-future computation might make this feasible (which is in line with trends – e.g., large-scale experiments with thousands of GPU-hours are becoming common).
- **Analysis of Emergent Systems:** When a system re-designs itself, it can become a “black box.” Understanding *why* the evolved architecture took a certain form, or *what* exactly the agent is imagining in the what-if engine, can be difficult. We observed interesting behaviors and could guess their reasons, but a more rigorous interpretability analysis would be needed to truly trust and understand such agents. This touches on AI safety as well – an agent driven by intrinsic goals might eventually pursue dangerous strategies if not properly constrained. Ensuring that some oversight or alignment is kept (perhaps by incorporating human feedback at meta-level) could be necessary if one were to deploy SPARSE-like agents in the real world.
- **No Objective = No Gauge of Progress?** In later stages we removed explicit objectives which makes it harder to measure progress in training. We relied on proxies like novelty or multi-environment success, but those are broad. One could argue the agent might meander or waste time in the absence of pressure. Indeed, purely open-ended systems like **Tierra** or other ALife simulations often create a lot of interesting patterns but not necessarily ones aligned with human notions of useful behavior. We tried to strike a balance by still having tasks (the environments) and the GENESIS oversight ensuring task performance. But as we give more freedom, there's a fine line between creative exploration and aimless wandering. Designing intrinsic objectives that are *just right* (not too restrictive, not too lax) is an art in itself.

Broader Impact and Future Work

SPARSE is a step toward what one might call **Artificial General Intelligence (AGI) incubators** – environments and training regimes that don't target a narrow skill, but cultivate general problem-solving ability and adaptability. By combining multiple mechanisms, we have created a richer training ground for an AI agent. Future work can extend this in many directions:

- **Scaling Up Environments:** Introduce not just 27 preset worlds, but an open-ended environment generator that continuously creates new worlds of increasing difficulty (inspired by POET ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#))). This would really push the open-ended aspect and see if our agent can keep scaling in competence.

- **Morphological Evolution:** We considered the neural architecture and policy, but one could also let the agent’s body evolve (if simulated robot) along with it. This would truly be a “brain-body co-evolution”. There’s evidence that co-evolving morphology and control yields innovative designs (). It would add another layer: Partition 5 perhaps, where the agent might modify its body or tools to better solve tasks. Combining that with what we have could birth very novel robotic lifeforms in simulation.
- **Real-world Transfer:** A practical angle is to use SPARSE to train agents (robots) in simulation that then transfer to reality. The diversity of environments and the use of an imagination module might make policies more robust to the unforeseen when deployed in the real world (they have effectively “seen it all” or can imagine new problems). Techniques like domain randomization already use environment diversity to achieve sim-to-real transfer; SPARSE could enhance this with an added dose of evolved robustness and meta-level adaptation. The what-if engine, if retained in the real robot, could help it avoid catastrophic trials by simulating them first on its internal model.
- **Neuroscience and Cognitive Science Insights:** Interestingly, our approach can also be used as a *model* for certain theories in cognitive science. For example, some theories suggest human cognition has multiple systems (fast intuitive experts vs slow deliberate reasoning – akin to our experts vs GENESIS – and a mental simulation faculty akin to imagination). SPARSE’s design loosely mirrors a “System 1 / System 2” hybrid. Studying how SPARSE agents behave might yield hypotheses about the benefits of having such dual systems and how they interact. Conversely, insights from cognitive science (like how humans balance exploration-exploitation, or how children play in imaginary scenarios) could inspire tweaks to SPARSE (like adding a play phase akin to Partition 4 but even when tasks are absent).
- **Philosophical Questions:** If an agent can simulate itself, adapt to new worlds, and generate novelty without explicit goals, one might ask: does it have a form of “will” or “intent”? Of course, it’s all within what we programmed, but as complexity grows, the boundary of control blurs. We set up the rules, but the outcomes might surpass our anticipation (indeed that’s the hope for emergent behaviors). There is a philosophical continuum from reactive agents to agents that we might call *creative*. SPARSE agents lie somewhere in between, and pushing further might lead to agents that we recognize as having autonomous creativity. This raises ethical considerations as well – an agent that is very general and introspective might need checks to ensure it doesn’t do unintended harm (much like any powerful AI).

In reflection, **the SPARSE protocol serves as both a research roadmap and a thought experiment.** Technically, it shows how integrating multiple AI paradigms can overcome each’s weaknesses: MoE alone doesn’t learn to gate itself, evolution alone can be inefficient, meta-learning alone can overfit to meta-training tasks, internal models alone don’t create new goals – but together, they form a more resilient whole. Philosophically, it invites us to consider that perhaps intelligence isn’t a single algorithm, but an ecology of processes. Evolution gave us our brain structure, learning refines our synapses, meta-cognition lets us think about thinking, and imagination lets us consider possibilities – SPARSE is an attempt to capture this multi-faceted evolution of intelligence in silico.

Conclusion

We have presented SPARSE, a Staged Partitioned Architecture Search protocol with Sub-Neural Dimensional Evaluation, as a comprehensive approach to training agents that are adaptable, general,

and capable of self-directed behavior. By progressing through staged training – from specialized experts to evolved architectures, to a meta-generalization coordination layer, and finally to an internal simulation engine – the SPARSE framework creates conditions for **emergent behaviors and structures** that no single method could produce in isolation. Our conceptual experiments across diverse simulated worlds illustrate the promise of this approach: agents that **learn how to learn**, repurpose their skills in novel contexts, and imagine new behaviors, all while largely free from the tether of manually designed rewards in later stages.

This work contributes to several research threads: it offers a new perspective on **mixture-of-experts**, showing how experts can be evolved and orchestrated rather than just fixed; it advances **neuroevolution** by demonstrating the synergy of evolution and learning in a staged manner; it resonates with **artificial life** by embracing open-ended, multi-environment evolution to yield unanticipated solutions; and it extends ideas of **model-based reinforcement learning** by using internal models not just for planning, but for creative exploration. In doing so, we hope to have sketched a pathway toward agents that possess a glimmer of **synthetic cognition** – not intelligent in the human sense, but autonomous in their growth and rich in their repertoire.

There are many avenues to build on this work. Empirically, implementing SPARSE at scale will be the next step, to validate these concepts with measurable outcomes and perhaps discover phenomena we didn't predict. The framework can be applied to various domains (robotics, game-playing, even intellectual tasks) to see how it generalizes. Each of the components (GENESIS layer, what-if engine, etc.) can be individually improved with the latest techniques (e.g., using transformer-based world models, or more principled meta-learning algorithms). We also foresee adapting SPARSE to continual learning settings, where an agent might face an endless sequence of tasks – SPARSE could naturally extend to that, given its open-ended design.

In conclusion, SPARSE represents a **holistic approach to agent development**. It does not yield just an agent that solves X or Y, but an agent that is an environment unto itself – one where evolutionary innovation, learning, and imagination interplay. By framing training as a multi-stage process that mirrors the stages of cognitive development (specialization, growth, integration, introspection), we move toward training AI systems that are not just problem solvers, but *problem creators and explorers*. We view this work as a step toward a future where AI agents are less like engineered tools and more like **artificial organisms**, exhibiting adaptive intelligence that grows and evolves over their lifetimes. The challenges are substantial, but the potential rewards – in terms of AI capability and understanding of cognition – are immense. SPARSE invites the community to rethink how we might **ignite an open-ended cognitive evolution** within machines, one that could eventually approach the open-endedness of natural evolution in generating intelligence ([\[1901.01753\] Paired Open-Ended Trailblazer \(POET\): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions](#)).